

# Home Mortgage Disclosure Act Data

Code ▾

Description: Cross-section data on the Home Mortgage Disclosure Act (HMDA). A data frame containing 2,380 observations on 14 variables.

Hide

```
library(AER)
data(HMDA)
dataset<- data.frame(HMDA)
names(dataset)
```

```
[1] "deny"      "pirat"     "hirat"     "lvrat"     "chist"
[6] "mhist"     "phist"     "unemp"     "selfemp"   "insurance"
[11] "condomin" "afam"      "single"    "hschool"
```

Hide

```
dim(dataset)
```

```
[1] 2380  14
```

Hide

```
str(dataset)
```

```
'data.frame':  2380 obs. of  14 variables:
 $ deny      : Factor w/ 2 levels "no","yes": 1 1 1 1 1 1 1 1 2 1 ...
 $ pirat     : num  0.221 0.265 0.372 0.32 0.36 ...
 $ hirat     : num  0.221 0.265 0.248 0.25 0.35 ...
 $ lvrat     : num  0.8 0.922 0.92 0.86 0.6 ...
 $ chist     : Factor w/ 6 levels "1","2","3","4",...: 5 2 1 1 1 1 1 2 2 2 ...
 $ mhist     : Factor w/ 4 levels "1","2","3","4": 2 2 2 2 1 1 2 2 2 1 ...
 $ phist     : Factor w/ 2 levels "no","yes": 1 1 1 1 1 1 1 1 1 1 ...
 $ unemp     : num  3.9 3.2 3.2 4.3 3.2 ...
 $ selfemp   : Factor w/ 2 levels "no","yes": 1 1 1 1 1 1 1 1 1 1 ...
 $ insurance: Factor w/ 2 levels "no","yes": 1 1 1 1 1 1 1 1 2 1 ...
 $ condominium : Factor w/ 2 levels "no","yes": 1 1 1 1 1 1 2 1 1 1 ...
 $ afam      : Factor w/ 2 levels "no","yes": 1 1 1 1 1 1 1 1 1 1 ...
 $ single    : Factor w/ 2 levels "no","yes": 1 2 1 1 1 1 2 1 1 2 ...
 $ hschool   : Factor w/ 2 levels "no","yes": 2 2 2 2 2 2 2 2 2 2 ...
```

Hide

```
summary(dataset)
```

```

deny      pirat      hirat      lvrat
no :2095   Min.   :0.0000   Min.   :0.0000   Min.   :0.0200
yes: 285   1st Qu.:0.2800   1st Qu.:0.2140   1st Qu.:0.6527
           Median :0.3300   Median :0.2600   Median :0.7795
           Mean   :0.3308   Mean   :0.2553   Mean   :0.7378
           3rd Qu.:0.3700   3rd Qu.:0.2988   3rd Qu.:0.8685
           Max.   :3.0000   Max.   :3.0000   Max.   :1.9500

chist     mhist     phist      unemp      selfemp
1:1353    1: 747    no :2205   Min.    : 1.800   no :2103
2: 441    2:1571   yes: 175   1st Qu.: 3.100   yes: 277
3: 126    3:  41           Median : 3.200
4:  77    4:  21           Mean  : 3.774
5: 182           3rd Qu.: 3.900
6: 201           Max.   :10.600

insurance  condominium  afam      single  hschool
no :2332   no :1694   no :2041   no :1444   no :  39
yes:  48   yes: 686   yes: 339   yes: 936   yes:2341

```

Hide

```
head(dataset)
```

deny <fctr>	pirat <dbl>	hirat <dbl>	lvrat <dbl>	chist <fctr>	mhist <fctr>	phist <fctr>	unemp <dbl>	selfemp <fctr>
1 no	0.221	0.221	0.8000000	5	2	no	3.9	no
2 no	0.265	0.265	0.9218750	2	2	no	3.2	no
3 no	0.372	0.248	0.9203980	1	2	no	3.2	no
4 no	0.320	0.250	0.8604651	1	2	no	4.3	no
5 no	0.360	0.350	0.6000000	1	1	no	3.2	no
6 no	0.240	0.170	0.5105263	1	1	no	3.9	no

6 rows | 1-10 of 14 columns

Hide

```
sapply(dataset,function(x) sum(is.na(x)))
```

```

deny      pirat      hirat      lvrat      chist      mhist
  0         0         0         0         0         0
phist     unemp     selfemp insurance  condominium  afam
  0         0         0         0         0         0
single    hschool
  0         0


```

As we can see data is clean, there are no missing values. Categorical values are already defined and correctly labeled. `pirat`, `hirat`, `lvrat`, `phist` are left skewed.

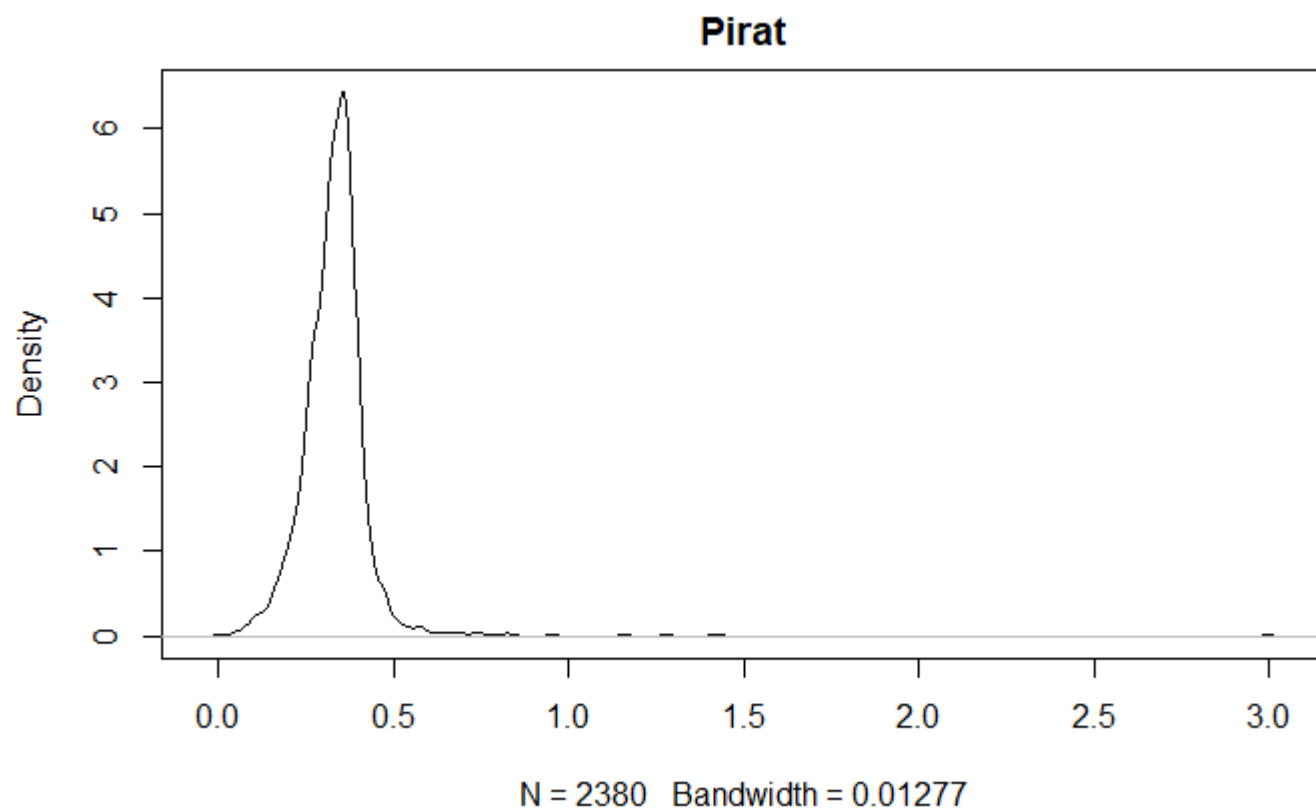
QQ-plot and density plots for payment to income ratio

[Hide](#)

```
qqnorm(dataset$pirat)
qqline(dataset$pirat,col='blue')
```

[Hide](#)

```
plot(density(dataset$pirat),main='Pirat')
```




The plot looks right skewed with few outliers.

QQ-plot and density plots housing expense to income ratio

Hide


```
qqnorm(dataset$hirat)
qqline(dataset$hirat,col='blue')
```

## Normal Q-Q Plot

[Hide](#)

```
plot(density(dataset$hirat),main='hirat')
```

## hirat




N = 2380 Bandwidth = 0.01203

The plot looks right skewed with few outliers.

## QQ-plot and density plots Loan to value ratio

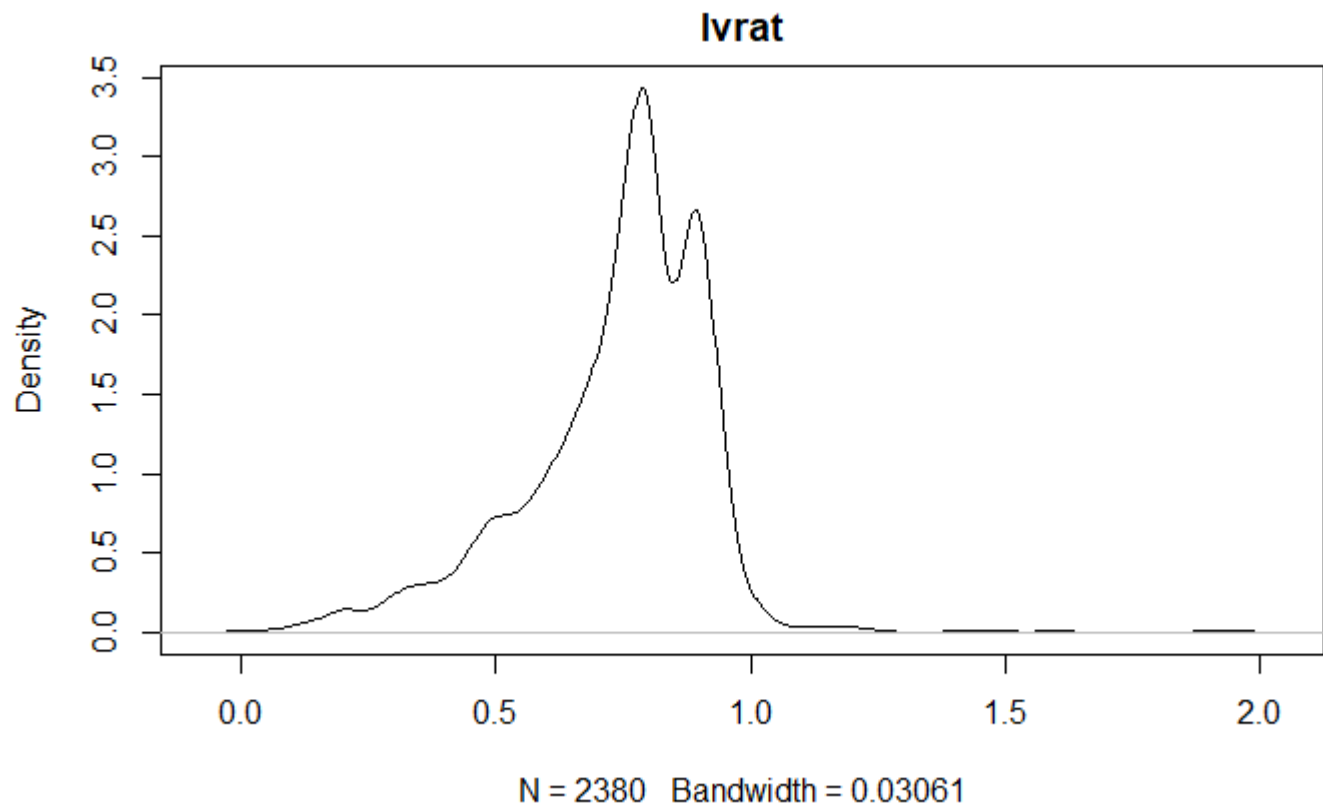
Hide

```
qqnorm(dataset$lvrat)  
qqline(dataset$lvrat,col='blue')
```



Hide

```
plot(density(dataset$lvrat),main='lvrat')
```




QQ-plot and density plots unemployment

Hide


```
qqnorm(dataset$unemp)
qqline(dataset$unemp,col='blue')
```

## Normal Q-Q Plot

[Hide](#)

```
plot(density(dataset$unemp),main='Unemp')
```

## Unemp




N = 2380 Bandwidth = 0.1135

## Box-Plots




Hide

```
library(ggplot2)
p10 <- ggplot(dataset, aes(x = deny, y = unemp)) +
  geom_boxplot()
p10
```



Hide

```
p11 <- ggplot(dataset, aes(x = deny, y = pirat)) +
  geom_boxplot()
p11
```



From the various QQ-Plots and Box-Plots we can conclude outliers are present.

Model fitting:

We split the data into two chunks: training and testing set. The training set will be used to fit our model which we will be testing over the testing set.

[Hide](#)

```
dt = sort(sample(nrow(dataset), nrow(dataset)*.8))
train<-dataset[dt,]
test<-dataset[-dt,]
```

Now, let's fit the model.

[Hide](#)

```
model1 <- glm(deny~.,family=binomial(link='logit'),data=train)
summary(model1)
```

```
Call:
glm(formula = deny ~ ., family = binomial(link = "logit"), data = train)
```

Deviance Residuals:

Min	1Q	Median	3Q	Max
-2.7851	-0.4209	-0.2986	-0.2052	3.0703

Coefficients:

	Estimate	Std. Error	z value	Pr(> z )	
(Intercept)	-6.06081	0.80635	-7.516	5.63e-14	***
pirat	5.68727	1.26257	4.505	6.65e-06	***
hirat	-0.98954	1.43543	-0.689	0.490590	
lvrat	2.16966	0.60116	3.609	0.000307	***
chist2	0.58326	0.23845	2.446	0.014445	*
chist3	0.65820	0.36733	1.792	0.073154	.
chist4	1.41414	0.39140	3.613	0.000303	***
chist5	1.14123	0.27684	4.122	3.75e-05	***
chist6	1.51969	0.25950	5.856	4.74e-09	***
mhist2	0.25891	0.22035	1.175	0.239992	
mhist3	0.03562	0.54847	0.065	0.948225	
mhist4	0.10461	0.85091	0.123	0.902152	
phistyes	1.42219	0.23797	5.976	2.28e-09	***
unemp	0.06084	0.03990	1.525	0.127331	
selfempyes	0.80705	0.24182	3.337	0.000846	***
insuranceyes	4.43250	0.56881	7.793	6.57e-15	***
condominyes	-0.07214	0.19511	-0.370	0.711564	
afamyas	0.52060	0.20872	2.494	0.012621	*
singleyes	0.33171	0.17982	1.845	0.065089	.
hschoolyes	-1.03454	0.48473	-2.134	0.032820	*

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 1363.10 on 1903 degrees of freedom

Residual deviance: 993.29 on 1884 degrees of freedom

AIC: 1033.3

Number of Fisher Scoring iterations: 6

By using function `summary()` we obtain the results of our model. Interpreting the results of our logistic regression model: Now we can analyze the fitting and interpret what the model is telling us. First of all, we can see that `hirat`, `mhist3`, `mhist4`, `unemp` and `condominyes` are not statistically significant. Whereas `pirat`, `phistyes`, `insuranceyes`, `afamyas` are statistically significant variables based on the p-values and AIC is 1024.9.

Now we can run the `anova()` function on the model to analyze the table of deviance

Hide

```
anova(model1, test="Chisq")
```

## Analysis of Deviance Table

Model: binomial, link: logit

Response: deny

Terms added sequentially (first to last)

	Df	Deviance	Resid. Df	Resid. Dev	Pr(>Chi)	
NULL			1903	1363.10		
pirat	1	62.871	1902	1300.23	2.207e-15	***
hirat	1	1.248	1901	1298.98	0.263860	
lvrat	1	49.191	1900	1249.79	2.322e-12	***
chist	5	84.696	1895	1165.09	< 2.2e-16	***
mhist	3	4.496	1892	1160.59	0.212618	
phist	1	34.695	1891	1125.90	3.856e-09	***
unemp	1	4.365	1890	1121.53	0.036691	*
selfemp	1	9.271	1889	1112.26	0.002329	**
insurance	1	105.444	1888	1006.82	< 2.2e-16	***
condomin	1	0.378	1887	1006.44	0.538540	
afam	1	6.488	1886	999.95	0.010859	*
single	1	2.595	1885	997.36	0.107196	
hschool	1	4.070	1884	993.29	0.043640	*

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

The difference between the null deviance and the residual deviance shows how our model is doing against the null model (a model with only the intercept). The wider this gap, the better. Analyzing the table we can see the drop in deviance when adding each variable one at a time. Again, adding lvrat, chist, phist, afam and hschool significantly reduces the residual deviance. A large p-value here indicates that the model without the variable explains more or less the same amount of variation. Ultimately what you would like to see is a significant drop in deviance and the AIC.

Hide

```
model2 <- glm(deny~pirat+lvrat+chist+phist+selfemp+insurance+afam+single+hschool,family=binomial
(link='logit'),data=train)
summary(model2)
```

Call:

```
glm(formula = deny ~ pirat + lvrat + chist + phist + selfemp +
     insurance + afam + single + hschool, family = binomial(link = "logit"),
     data = train)
```

Deviance Residuals:

Min	1Q	Median	3Q	Max
-2.9276	-0.4199	-0.3009	-0.2154	3.1311

Coefficients:

	Estimate	Std. Error	z value	Pr(> z )	
(Intercept)	-5.6749	0.7472	-7.595	3.07e-14	***
pirat	5.0779	0.9352	5.430	5.65e-08	***
lvrat	2.2599	0.5949	3.799	0.000146	***
chist2	0.5991	0.2366	2.532	0.011349	*
chist3	0.6451	0.3628	1.778	0.075351	.
chist4	1.3837	0.3882	3.564	0.000365	***
chist5	1.1493	0.2756	4.170	3.05e-05	***
chist6	1.5505	0.2544	6.096	1.09e-09	***
phistyes	1.4151	0.2373	5.963	2.47e-09	***
selfempyes	0.8260	0.2370	3.485	0.000492	***
insuranceyes	4.4635	0.5661	7.884	3.16e-15	***
afamyes	0.5115	0.2048	2.498	0.012494	*
singleyes	0.3339	0.1721	1.940	0.052357	.
hschoolyes	-1.1515	0.4782	-2.408	0.016028	*

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 1363.10 on 1903 degrees of freedom  
 Residual deviance: 997.53 on 1890 degrees of freedom  
 AIC: 1025.5

Number of Fisher Scoring iterations: 6

In model two hirat,unemp,mhist and condomin are removed. AIC is 1016.6.

Hide

```
anova(model2, test="Chisq")
```

## Analysis of Deviance Table

Model: binomial, link: logit

Response: deny

Terms added sequentially (first to last)

	Df	Deviance	Resid. Df	Resid. Dev	Pr(>Chi)	
NULL			1903	1363.10		
pirat	1	62.871	1902	1300.23	2.207e-15	***
lvrat	1	48.819	1901	1251.41	2.807e-12	***
chist	5	86.015	1896	1165.39	< 2.2e-16	***
phist	1	34.562	1895	1130.83	4.128e-09	***
selfemp	1	9.172	1894	1121.66	0.002458	**
insurance	1	108.671	1893	1012.99	< 2.2e-16	***
afam	1	7.347	1892	1005.64	0.006716	**
single	1	3.051	1891	1002.59	0.080708	.
hschool	1	5.057	1890	997.53	0.024533	*

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

[Hide](#)

```
model3 <- glm(deny~pirat+lvrat+chist+phist+insurance+afam+single,family=binomial(link='logit'),data=train)
summary(model3)
```

Call:

```
glm(formula = deny ~ pirat + lvrat + chist + phist + insurance +
     afam + single, family = binomial(link = "logit"), data = train)
```

Deviance Residuals:

Min	1Q	Median	3Q	Max
-2.9383	-0.4263	-0.3119	-0.2219	3.0687

Coefficients:

	Estimate	Std. Error	z value	Pr(> z )	
(Intercept)	-6.5964	0.5863	-11.252	< 2e-16	***
pirat	5.0623	0.9326	5.428	5.70e-08	***
lvrat	2.1663	0.5957	3.636	0.000277	***
chist2	0.6378	0.2349	2.715	0.006637	**
chist3	0.6854	0.3578	1.916	0.055385	.
chist4	1.4107	0.3803	3.710	0.000207	***
chist5	1.2025	0.2736	4.395	1.11e-05	***
chist6	1.5011	0.2524	5.947	2.73e-09	***
phistyes	1.4213	0.2358	6.027	1.67e-09	***
insuranceyes	4.4289	0.5666	7.816	5.45e-15	***
afamyas	0.4951	0.2030	2.439	0.014717	*
singleyes	0.2901	0.1698	1.708	0.087626	.

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 1363.1 on 1903 degrees of freedom  
 Residual deviance: 1013.1 on 1892 degrees of freedom  
 AIC: 1037.1

Number of Fisher Scoring iterations: 6

Hide

```
anova(model3, test="Chisq")
```

## Analysis of Deviance Table

Model: binomial, link: logit

Response: deny

Terms added sequentially (first to last)

	Df	Deviance	Resid. Df	Resid. Dev	Pr(>Chi)
NULL			1903	1363.1	
pirat	1	62.871	1902	1300.2	2.207e-15 ***
lvrat	1	48.819	1901	1251.4	2.807e-12 ***
chist	5	86.015	1896	1165.4	< 2.2e-16 ***
phist	1	34.562	1895	1130.8	4.128e-09 ***
insurance	1	108.191	1894	1022.6	< 2.2e-16 ***
afam	1	6.590	1893	1016.0	0.01026 *
single	1	2.902	1892	1013.1	0.08846 .

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Even we can see the p-value associated with selfemp and unemp is not significant as its large but removing the element from model increases the AIC value. So model3 is not a good model.

From here we conclude that model2 is the best. Now we will use forward selection to verify our model.

Hide

```
fit1<- glm(deny~pirat,family=binomial(link='logit'),data=train)
summary(fit1)
```



```
Call:
glm(formula = deny ~ pirat, family = binomial(link = "logit"),
     data = train)

Deviance Residuals:
    Min       1Q   Median       3Q      Max
-1.9669  -0.5198  -0.4617  -0.3756   2.8163

Coefficients:
              Estimate Std. Error z value Pr(>|z|)
(Intercept)  -4.1984     0.3156 -13.304  < 2e-16 ***
pirat         6.2911     0.8644   7.278 3.39e-13 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

    Null deviance: 1363.1  on 1903  degrees of freedom
Residual deviance: 1300.2  on 1902  degrees of freedom
AIC: 1304.2

Number of Fisher Scoring iterations: 5
```

[Hide](#)

```
fit2<-glm(deny~pirat+hirat,family=binomial(link='logit'),data=train)
summary(fit2)
```

Call:

```
glm(formula = deny ~ pirat + hirat, family = binomial(link = "logit"),
    data = train)
```

Deviance Residuals:

Min	1Q	Median	3Q	Max
-2.0510	-0.5246	-0.4557	-0.3771	2.7859

Coefficients:

	Estimate	Std. Error	z value	Pr(> z )
(Intercept)	-4.1006	0.3254	-12.603	< 2e-16 ***
pirat	6.9987	1.0654	6.569	5.06e-11 ***
hirat	-1.3066	1.1607	-1.126	0.26

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 1363.1 on 1903 degrees of freedom

Residual deviance: 1299.0 on 1901 degrees of freedom

AIC: 1305

Number of Fisher Scoring iterations: 5

Hide

```
fit3<-glm(deny~pirat+lvrat,family=binomial(link='logit'),data=train)
summary(fit3)
```

Call:

```
glm(formula = deny ~ pirat + lvrat, family = binomial(link = "logit"),
     data = train)
```

Deviance Residuals:

Min	1Q	Median	3Q	Max
-1.7020	-0.5453	-0.4289	-0.2914	3.1713

Coefficients:

	Estimate	Std. Error	z value	Pr(> z )
(Intercept)	-6.7419	0.5336	-12.635	< 2e-16 ***
pirat	6.0314	0.8813	6.844	7.70e-12 ***
lvrat	3.3823	0.5204	6.500	8.04e-11 ***

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 1363.1 on 1903 degrees of freedom

Residual deviance: 1251.4 on 1901 degrees of freedom

AIC: 1257.4

Number of Fisher Scoring iterations: 5

Hide

```
fit4<-glm(deny~pirat+lvrat+chist,family=binomial(link='logit'),data=train)
summary(fit4)
```

Call:

```
glm(formula = deny ~ pirat + lvrat + chist, family = binomial(link = "logit"),
    data = train)
```

Deviance Residuals:

Min	1Q	Median	3Q	Max
-1.9901	-0.4863	-0.3655	-0.2411	3.3151

Coefficients:

	Estimate	Std. Error	z value	Pr(> z )
(Intercept)	-7.0869	0.5529	-12.817	< 2e-16 ***
pirat	5.5324	0.9030	6.127	8.96e-10 ***
lvrat	3.2008	0.5350	5.983	2.19e-09 ***
chist2	0.6752	0.2134	3.164	0.001555 **
chist3	0.9876	0.3140	3.145	0.001660 **
chist4	1.2930	0.3697	3.498	0.000469 ***
chist5	1.2534	0.2489	5.035	4.77e-07 ***
chist6	1.9118	0.2175	8.788	< 2e-16 ***

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 1363.1 on 1903 degrees of freedom  
 Residual deviance: 1165.4 on 1896 degrees of freedom  
 AIC: 1181.4

Number of Fisher Scoring iterations: 6

Hide

```
fit5<-glm(deny~pirat+lvrat+chist+mhist,family=binomial(link='logit'),data=train)
summary(fit5)
```

Call:

```
glm(formula = deny ~ pirat + lvrat + chist + mhist, family = binomial(link = "logit"),
    data = train)
```

Deviance Residuals:

Min	1Q	Median	3Q	Max
-1.9658	-0.4863	-0.3642	-0.2350	3.2529

Coefficients:

	Estimate	Std. Error	z value	Pr(> z )	
(Intercept)	-7.2324	0.5614	-12.882	< 2e-16	***
pirat	5.5719	0.9136	6.099	1.07e-09	***
lvrat	3.0173	0.5398	5.590	2.27e-08	***
chist2	0.6913	0.2137	3.234	0.001219	**
chist3	1.0103	0.3180	3.177	0.001488	**
chist4	1.2378	0.3710	3.337	0.000848	***
chist5	1.2469	0.2491	5.006	5.55e-07	***
chist6	1.8603	0.2206	8.432	< 2e-16	***
mhist2	0.3782	0.1949	1.941	0.052266	.
mhist3	0.2038	0.5271	0.387	0.698992	
mhist4	0.4431	0.7276	0.609	0.542584	

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 1363.1 on 1903 degrees of freedom  
 Residual deviance: 1161.4 on 1893 degrees of freedom  
 AIC: 1183.4

Number of Fisher Scoring iterations: 6

Hide

```
fit6<-glm(deny~pirat+lvrat+chist+phist,family=binomial(link='logit'),data=train)
summary(fit6)
```

Call:

```
glm(formula = deny ~ pirat + lvrat + chist + phist, family = binomial(link = "logit"),
    data = train)
```

Deviance Residuals:

Min	1Q	Median	3Q	Max
-1.7269	-0.4704	-0.3540	-0.2379	3.3042

Coefficients:

	Estimate	Std. Error	z value	Pr(> z )	
(Intercept)	-6.9880	0.5574	-12.538	< 2e-16	***
pirat	5.2821	0.9100	5.804	6.46e-09	***
lvrat	3.1080	0.5388	5.768	8.01e-09	***
chist2	0.6188	0.2157	2.868	0.00413	**
chist3	0.7358	0.3254	2.261	0.02377	*
chist4	1.3161	0.3711	3.547	0.00039	***
chist5	1.1144	0.2554	4.363	1.29e-05	***
chist6	1.5057	0.2345	6.420	1.36e-10	***
phistyes	1.3638	0.2248	6.066	1.31e-09	***

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 1363.1 on 1903 degrees of freedom  
 Residual deviance: 1130.8 on 1895 degrees of freedom  
 AIC: 1148.8

Number of Fisher Scoring iterations: 6

Hide

```
fit7<-glm(deny~pirat+lvrat+chist+phist+unemp,family=binomial(link='logit'),data=train)
summary(fit7)
```

Call:

```
glm(formula = deny ~ pirat + lvrat + chist + phist + unemp, family = binomial(link = "logit"),
    data = train)
```

Deviance Residuals:

Min	1Q	Median	3Q	Max
-1.9589	-0.4668	-0.3497	-0.2328	3.3526

Coefficients:

	Estimate	Std. Error	z value	Pr(> z )
(Intercept)	-7.30941	0.58100	-12.581	< 2e-16 ***
pirat	5.26079	0.91344	5.759	8.44e-09 ***
lvrat	3.14372	0.53943	5.828	5.61e-09 ***
chist2	0.59518	0.21641	2.750	0.005954 **
chist3	0.72200	0.32565	2.217	0.026618 *
chist4	1.34082	0.37164	3.608	0.000309 ***
chist5	1.11903	0.25611	4.369	1.25e-05 ***
chist6	1.52180	0.23495	6.477	9.35e-11 ***
phistyes	1.36543	0.22451	6.082	1.19e-09 ***
unemp	0.07829	0.03580	2.187	0.028768 *

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 1363.1 on 1903 degrees of freedom

Residual deviance: 1126.4 on 1894 degrees of freedom

AIC: 1146.4

Number of Fisher Scoring iterations: 6

Hide

```
fit8<-glm(deny~pirat+lvrat+chist+phist+selfemp,family=binomial(link='logit'),data=train)
summary(fit8)
```

Call:

```
glm(formula = deny ~ pirat + lvrat + chist + phist + selfemp,
     family = binomial(link = "logit"), data = train)
```

Deviance Residuals:

Min	1Q	Median	3Q	Max
-2.0211	-0.4652	-0.3503	-0.2340	3.3500

Coefficients:

	Estimate	Std. Error	z value	Pr(> z )	
(Intercept)	-7.1662	0.5619	-12.752	< 2e-16	***
pirat	5.3239	0.9122	5.836	5.35e-09	***
lvrat	3.2009	0.5427	5.899	3.66e-09	***
chist2	0.6025	0.2166	2.782	0.005405	**
chist3	0.7264	0.3261	2.227	0.025921	*
chist4	1.3506	0.3733	3.618	0.000297	***
chist5	1.0844	0.2576	4.210	2.56e-05	***
chist6	1.5488	0.2357	6.572	4.96e-11	***
phistyes	1.3602	0.2248	6.051	1.44e-09	***
selfempyes	0.7121	0.2256	3.156	0.001597	**

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 1363.1 on 1903 degrees of freedom

Residual deviance: 1121.7 on 1894 degrees of freedom

AIC: 1141.7

Number of Fisher Scoring iterations: 6

Hide

```
fit9<-glm(deny~pirat+lvrat+chist+phist+insurance,family=binomial(link='logit'),data=train)
summary(fit9)
```



Call:

```
glm(formula = deny ~ pirat + lvrat + chist + phist + insurance,
     family = binomial(link = "logit"), data = train)
```

Deviance Residuals:

Min	1Q	Median	3Q	Max
-2.9240	-0.4257	-0.3192	-0.2250	3.2473

Coefficients:

	Estimate	Std. Error	z value	Pr(> z )	
(Intercept)	-6.6588	0.5818	-11.445	< 2e-16	***
pirat	5.2216	0.9366	5.575	2.47e-08	***
lvrat	2.4085	0.5871	4.103	4.09e-05	***
chist2	0.6490	0.2342	2.771	0.00559	**
chist3	0.7494	0.3534	2.121	0.03395	*
chist4	1.5678	0.3752	4.179	2.93e-05	***
chist5	1.2566	0.2693	4.666	3.08e-06	***
chist6	1.6175	0.2476	6.533	6.44e-11	***
phistyes	1.4524	0.2323	6.252	4.05e-10	***
insuranceyes	4.4254	0.5653	7.828	4.95e-15	***

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 1363.1 on 1903 degrees of freedom

Residual deviance: 1022.6 on 1894 degrees of freedom

AIC: 1042.6

Number of Fisher Scoring iterations: 6

Hide

```
anova(fit9, test="Chisq")
```

## Analysis of Deviance Table

Model: binomial, link: logit

Response: deny

Terms added sequentially (first to last)

	Df	Deviance	Resid. Df	Resid. Dev	Pr(>Chi)
NULL			1903	1363.1	
pirat	1	62.871	1902	1300.2	2.207e-15 ***
lvrat	1	48.819	1901	1251.4	2.807e-12 ***
chist	5	86.015	1896	1165.4	< 2.2e-16 ***
phist	1	34.562	1895	1130.8	4.128e-09 ***
insurance	1	108.191	1894	1022.6	< 2.2e-16 ***

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Hide

```
fit10<-glm(deny~pirat+lvrat+chist+phist+insurance+condomin,family=binomial(link='logit'),data=train)
summary(fit10)
```

Call:

```
glm(formula = deny ~ pirat + lvrat + chist + phist + insurance +
     condomin, family = binomial(link = "logit"), data = train)
```

Deviance Residuals:

Min	1Q	Median	3Q	Max
-2.9458	-0.4241	-0.3190	-0.2261	3.2216

Coefficients:

	Estimate	Std. Error	z value	Pr(> z )	
(Intercept)	-6.6743	0.5819	-11.471	< 2e-16	***
pirat	5.2246	0.9353	5.586	2.32e-08	***
lvrat	2.3880	0.5877	4.063	4.85e-05	***
chist2	0.6573	0.2348	2.799	0.00512	**
chist3	0.7495	0.3532	2.122	0.03386	*
chist4	1.5438	0.3774	4.090	4.30e-05	***
chist5	1.2546	0.2694	4.656	3.22e-06	***
chist6	1.6093	0.2481	6.488	8.72e-11	***
phistyes	1.4519	0.2324	6.247	4.17e-10	***
insuranceyes	4.4344	0.5658	7.837	4.60e-15	***
condominyes	0.1023	0.1813	0.564	0.57250	

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 1363.1 on 1903 degrees of freedom  
 Residual deviance: 1022.3 on 1893 degrees of freedom  
 AIC: 1044.3

Number of Fisher Scoring iterations: 6

Hide

```
fit11<-glm(deny~pirat+lvrat+chist+phist+insurance+afam,family=binomial(link='logit'),data=train)
summary(fit11)
```

Call:

```
glm(formula = deny ~ pirat + lvrat + chist + phist + insurance +
     afam, family = binomial(link = "logit"), data = train)
```

Deviance Residuals:

Min	1Q	Median	3Q	Max
-2.9968	-0.4263	-0.3163	-0.2259	3.1132

Coefficients:

	Estimate	Std. Error	z value	Pr(> z )	
(Intercept)	-6.4904	0.5813	-11.165	< 2e-16	***
pirat	5.0777	0.9348	5.432	5.57e-08	***
lvrat	2.1739	0.5937	3.662	0.000251	***
chist2	0.6372	0.2345	2.717	0.006593	**
chist3	0.7061	0.3550	1.989	0.046708	*
chist4	1.4598	0.3790	3.852	0.000117	***
chist5	1.1877	0.2727	4.355	1.33e-05	***
chist6	1.5071	0.2525	5.968	2.40e-09	***
phistyes	1.4046	0.2351	5.975	2.31e-09	***
insuranceyes	4.4462	0.5659	7.857	3.95e-15	***
afamyes	0.5299	0.2020	2.624	0.008696	**

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 1363.1 on 1903 degrees of freedom  
 Residual deviance: 1016.0 on 1893 degrees of freedom  
 AIC: 1038

Number of Fisher Scoring iterations: 6

Hide

```
fit12<-glm(deny~pirat+lvrat+chist+phist+insurance+afam+single,family=binomial(link='logit'),data
=train)
summary(fit12)
```

Call:

```
glm(formula = deny ~ pirat + lvrat + chist + phist + insurance +
     afam + single, family = binomial(link = "logit"), data = train)
```

Deviance Residuals:

Min	1Q	Median	3Q	Max
-2.9383	-0.4263	-0.3119	-0.2219	3.0687

Coefficients:

	Estimate	Std. Error	z value	Pr(> z )	
(Intercept)	-6.5964	0.5863	-11.252	< 2e-16	***
pirat	5.0623	0.9326	5.428	5.70e-08	***
lvrat	2.1663	0.5957	3.636	0.000277	***
chist2	0.6378	0.2349	2.715	0.006637	**
chist3	0.6854	0.3578	1.916	0.055385	.
chist4	1.4107	0.3803	3.710	0.000207	***
chist5	1.2025	0.2736	4.395	1.11e-05	***
chist6	1.5011	0.2524	5.947	2.73e-09	***
phistyes	1.4213	0.2358	6.027	1.67e-09	***
insuranceyes	4.4289	0.5666	7.816	5.45e-15	***
afamyas	0.4951	0.2030	2.439	0.014717	*
singleyes	0.2901	0.1698	1.708	0.087626	.

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 1363.1 on 1903 degrees of freedom  
 Residual deviance: 1013.1 on 1892 degrees of freedom  
 AIC: 1037.1

Number of Fisher Scoring iterations: 6

Hide

```
fit13<-glm(deny~pirat+lvrat+chist+insurance+afam+single+hschool,family=binomial(link='logit'),da
ta=train)
summary(fit13)
```

Call:

```
glm(formula = deny ~ pirat + lvrat + chist + insurance + afam +
     single + hschool, family = binomial(link = "logit"), data = train)
```

Deviance Residuals:

Min	1Q	Median	3Q	Max
-3.0201	-0.4428	-0.3184	-0.2232	3.0774

Coefficients:

	Estimate	Std. Error	z value	Pr(> z )	
(Intercept)	-5.6056	0.7376	-7.600	2.96e-14	***
pirat	5.2335	0.9153	5.718	1.08e-08	***
lvrat	2.2710	0.5908	3.844	0.000121	***
chist2	0.6837	0.2329	2.936	0.003325	**
chist3	0.9553	0.3444	2.774	0.005541	**
chist4	1.3085	0.3821	3.424	0.000617	***
chist5	1.3324	0.2654	5.020	5.17e-07	***
chist6	1.9027	0.2363	8.051	8.24e-16	***
insuranceyes	4.4061	0.5669	7.772	7.75e-15	***
afamyes	0.5784	0.1966	2.941	0.003267	**
singleyes	0.2797	0.1674	1.671	0.094644	.
hschoolyes	-1.0924	0.4687	-2.331	0.019770	*

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 1363.1 on 1903 degrees of freedom  
 Residual deviance: 1042.3 on 1892 degrees of freedom  
 AIC: 1066.3

Number of Fisher Scoring iterations: 6

fit2 discard as AIC increases and p-value is too large fit3 is good AIC reduces fit4

fit5 even though the p-values are too large the AIC of the model decreases\*\* fit6 is good fit7 although the AIC remains same we can see the p-value associated with unemp is large so fit7 is not a good model fit8 AIC remains same, the p-value is greater than 0.05 so fit8 is discarded fit9 Even though the AIC increases, but the deviance decreases significantly and the associated p-value is significant we will keep fit9 fit10 discard based on p-values fit11 is good reduces AIC a lot fit12 is good fit13 discard

So fit12 is the best model. Comparing it with our previous model3

Hide

```
model3 <- glm(deny~pirat+lvrat+chist+phist+insurance+afam+single,family=binomial(link='logit'),data=train)
summary(model3)
```

Call:

```
glm(formula = deny ~ pirat + lvrat + chist + phist + insurance +
     afam + single, family = binomial(link = "logit"), data = train)
```

Deviance Residuals:

Min	1Q	Median	3Q	Max
-2.9383	-0.4263	-0.3119	-0.2219	3.0687

Coefficients:

	Estimate	Std. Error	z value	Pr(> z )	
(Intercept)	-6.5964	0.5863	-11.252	< 2e-16	***
pirat	5.0623	0.9326	5.428	5.70e-08	***
lvrat	2.1663	0.5957	3.636	0.000277	***
chist2	0.6378	0.2349	2.715	0.006637	**
chist3	0.6854	0.3578	1.916	0.055385	.
chist4	1.4107	0.3803	3.710	0.000207	***
chist5	1.2025	0.2736	4.395	1.11e-05	***
chist6	1.5011	0.2524	5.947	2.73e-09	***
phistyes	1.4213	0.2358	6.027	1.67e-09	***
insuranceyes	4.4289	0.5666	7.816	5.45e-15	***
afamyas	0.4951	0.2030	2.439	0.014717	*
singleyes	0.2901	0.1698	1.708	0.087626	.

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 1363.1 on 1903 degrees of freedom  
 Residual deviance: 1013.1 on 1892 degrees of freedom  
 AIC: 1037.1

Number of Fisher Scoring iterations: 6

Hide

```
model2 <- glm(deny~pirat+lvrat+chist+phist+selfemp+insurance+afam+single+hschool,family=binomial
(link='logit'),data=train)
fit12<-glm(deny~pirat+lvrat+chist+phist+insurance+afam+single,family=binomial(link='logit'),data
=train)
summary(fit12)
```

Call:

```
glm(formula = deny ~ pirat + lvrat + chist + phist + insurance +
     afam + single, family = binomial(link = "logit"), data = train)
```

Deviance Residuals:

Min	1Q	Median	3Q	Max
-2.9383	-0.4263	-0.3119	-0.2219	3.0687

Coefficients:

	Estimate	Std. Error	z value	Pr(> z )	
(Intercept)	-6.5964	0.5863	-11.252	< 2e-16	***
pirat	5.0623	0.9326	5.428	5.70e-08	***
lvrat	2.1663	0.5957	3.636	0.000277	***
chist2	0.6378	0.2349	2.715	0.006637	**
chist3	0.6854	0.3578	1.916	0.055385	.
chist4	1.4107	0.3803	3.710	0.000207	***
chist5	1.2025	0.2736	4.395	1.11e-05	***
chist6	1.5011	0.2524	5.947	2.73e-09	***
phistyes	1.4213	0.2358	6.027	1.67e-09	***
insuranceyes	4.4289	0.5666	7.816	5.45e-15	***
afamyas	0.4951	0.2030	2.439	0.014717	*
singleyes	0.2901	0.1698	1.708	0.087626	.

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 1363.1 on 1903 degrees of freedom  
 Residual deviance: 1013.1 on 1892 degrees of freedom  
 AIC: 1037.1

Number of Fisher Scoring iterations: 6

So from the forward model selection method we can conclude Model3 is the best.

So our model is : model3 <-

```
glm(deny~pirat+lvrat+chist+phist+insurance+afam+single,family=binomial(link='logit'),data=train)
```

Hide

```
fm1 <- lm(I(as.numeric(deny) - 1) ~ pirat+lvrat+chist+phist+insurance+afam+single, data = dataset)
summary(fm1)
```



Call:

```
lm(formula = I(as.numeric(deny) - 1) ~ pirat + lvrat + chist +  
    phist + insurance + afam + single, data = dataset)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-0.93240	-0.11506	-0.05449	-0.00940	1.08223

Coefficients:

	Estimate	Std. Error	t value	Pr(> t )	
(Intercept)	-0.18538	0.02903	-6.386	2.04e-10	***
pirat	0.43341	0.05461	7.937	3.16e-15	***
lvrat	0.09423	0.03327	2.832	0.00466	**
chist2	0.04236	0.01541	2.750	0.00601	**
chist3	0.05243	0.02643	1.984	0.04738	*
chist4	0.13951	0.03316	4.208	2.67e-05	***
chist5	0.11345	0.02246	5.051	4.74e-07	***
chist6	0.16169	0.02255	7.170	9.95e-13	***
phistyes	0.20813	0.02357	8.829	< 2e-16	***
insuranceyes	0.71292	0.04143	17.207	< 2e-16	***
afamyes	0.08340	0.01728	4.826	1.48e-06	***
singleyes	0.03355	0.01188	2.823	0.00479	**

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.2806 on 2368 degrees of freedom

Multiple R-squared: 0.2567, Adjusted R-squared: 0.2533

F-statistic: 74.36 on 11 and 2368 DF, p-value: < 2.2e-16

Hide

```
fm2 <- lm(I(as.numeric(deny) - 1) ~ pirat+lvrat+chist+insurance+afam+single+hschool, data = data  
set)  
summary(fm2)
```

Call:

```
lm(formula = I(as.numeric(deny) - 1) ~ pirat + lvrat + chist +
    insurance + afam + single + hschool, data = dataset)
```

Residuals:

Min	1Q	Median	3Q	Max
-1.01018	-0.12217	-0.05868	-0.00999	1.08383

Coefficients:

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	-0.07432	0.05553	-1.338	0.18096
pirat	0.45315	0.05552	8.163	5.27e-16 ***
lvrat	0.10746	0.03373	3.186	0.00146 **
chist2	0.04638	0.01563	2.968	0.00303 **
chist3	0.07994	0.02664	3.001	0.00272 **
chist4	0.13161	0.03369	3.906	9.64e-05 ***
chist5	0.13060	0.02270	5.753	9.87e-09 ***
chist6	0.21777	0.02200	9.900	< 2e-16 ***
insuranceyes	0.73613	0.04199	17.531	< 2e-16 ***
afamyes	0.09489	0.01749	5.427	6.33e-08 ***
singleyes	0.03379	0.01208	2.797	0.00519 **
hschoolyes	-0.12428	0.04641	-2.678	0.00746 **

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.2848 on 2368 degrees of freedom

Multiple R-squared: 0.2346, Adjusted R-squared: 0.231

F-statistic: 65.98 on 11 and 2368 DF, p-value: < 2.2e-16

Hide

```
fm3 <- lm(I(as.numeric(deny) - 1) ~., data = dataset)
summary(fm3)
```

Call:

```
lm(formula = I(as.numeric(deny) - 1) ~ ., data = dataset)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-0.92559	-0.12214	-0.05246	-0.00258	1.07741

Coefficients:

	Estimate	Std. Error	t value	Pr(> t )	
(Intercept)	-0.095597	0.057354	-1.667	0.09569	.
pirat	0.476759	0.088653	5.378	8.28e-08	***
hirat	-0.099505	0.097994	-1.015	0.31001	
lvrat	0.094726	0.033942	2.791	0.00530	**
chist2	0.038143	0.015483	2.464	0.01383	*
chist3	0.049182	0.026661	1.845	0.06520	.
chist4	0.135810	0.033372	4.070	4.86e-05	***
chist5	0.106922	0.022466	4.759	2.06e-06	***
chist6	0.160628	0.022773	7.054	2.28e-12	***
mhist2	0.018115	0.013432	1.349	0.17758	
mhist3	0.034372	0.045616	0.754	0.45122	
mhist4	0.027747	0.062691	0.443	0.65810	
phistyes	0.204367	0.023531	8.685	< 2e-16	***
unemp	0.004688	0.002903	1.615	0.10650	
selfempyes	0.056660	0.018305	3.095	0.00199	**
insuranceyes	0.713226	0.041360	17.244	< 2e-16	***
condominyes	-0.005913	0.013654	-0.433	0.66501	
afamyyes	0.084846	0.017537	4.838	1.40e-06	***
singleyes	0.035722	0.012599	2.835	0.00462	**
hschoolyes	-0.115475	0.045984	-2.511	0.01210	*

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

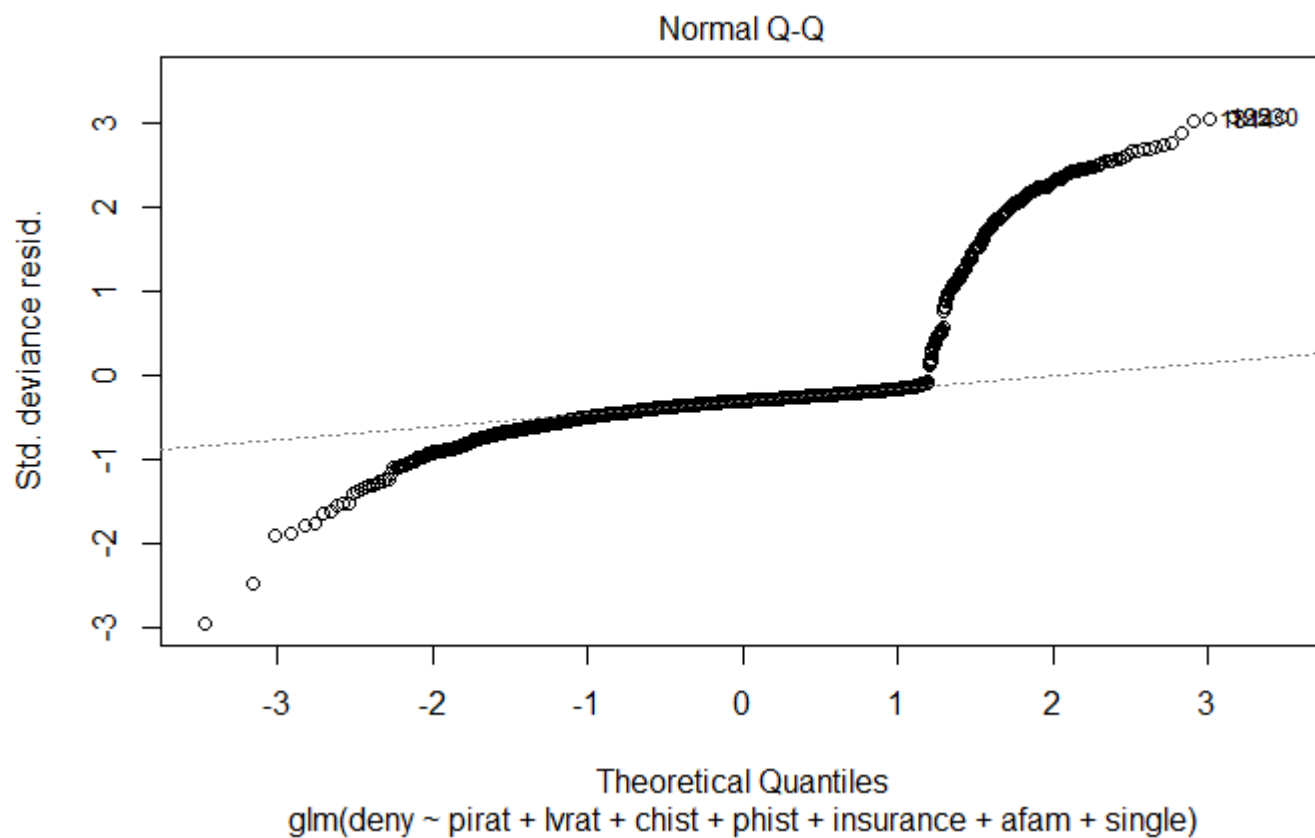
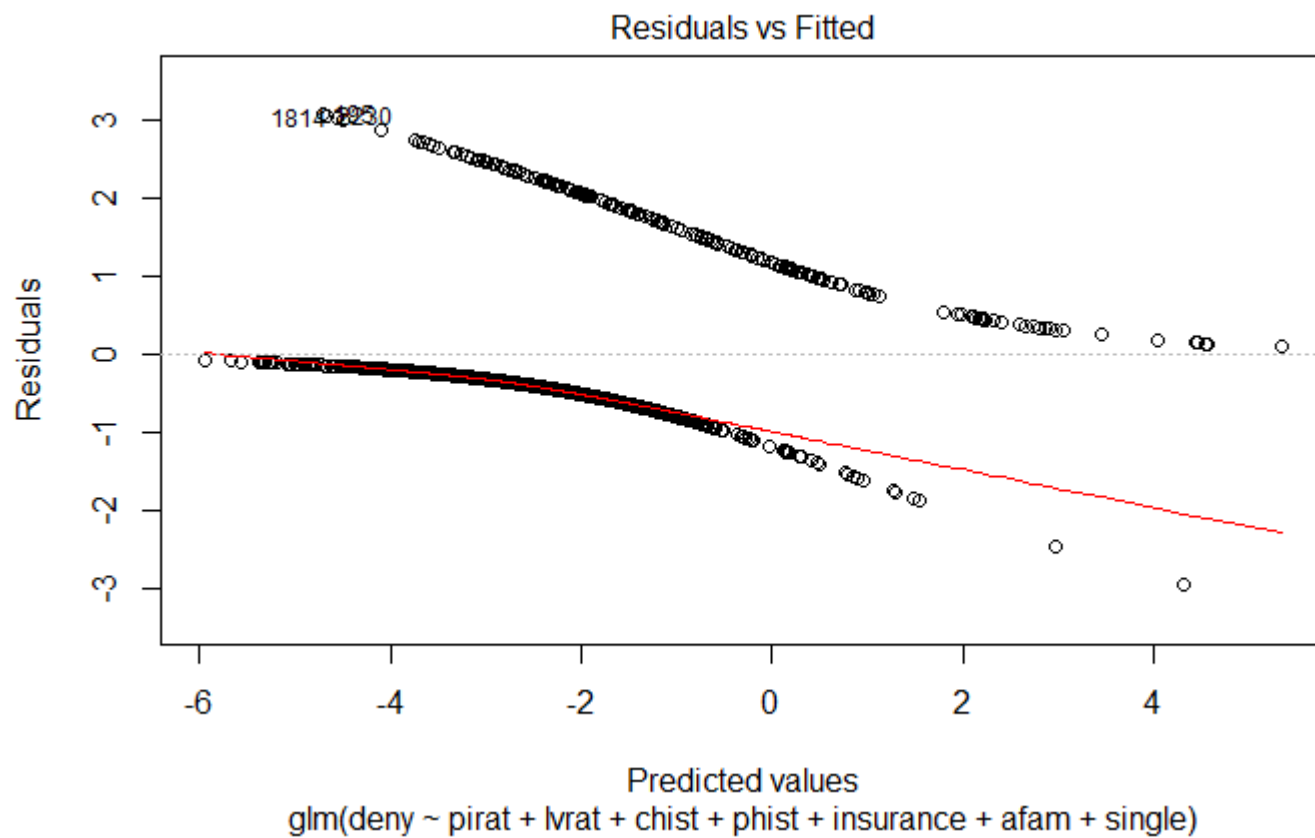
Residual standard error: 0.2796 on 2360 degrees of freedom


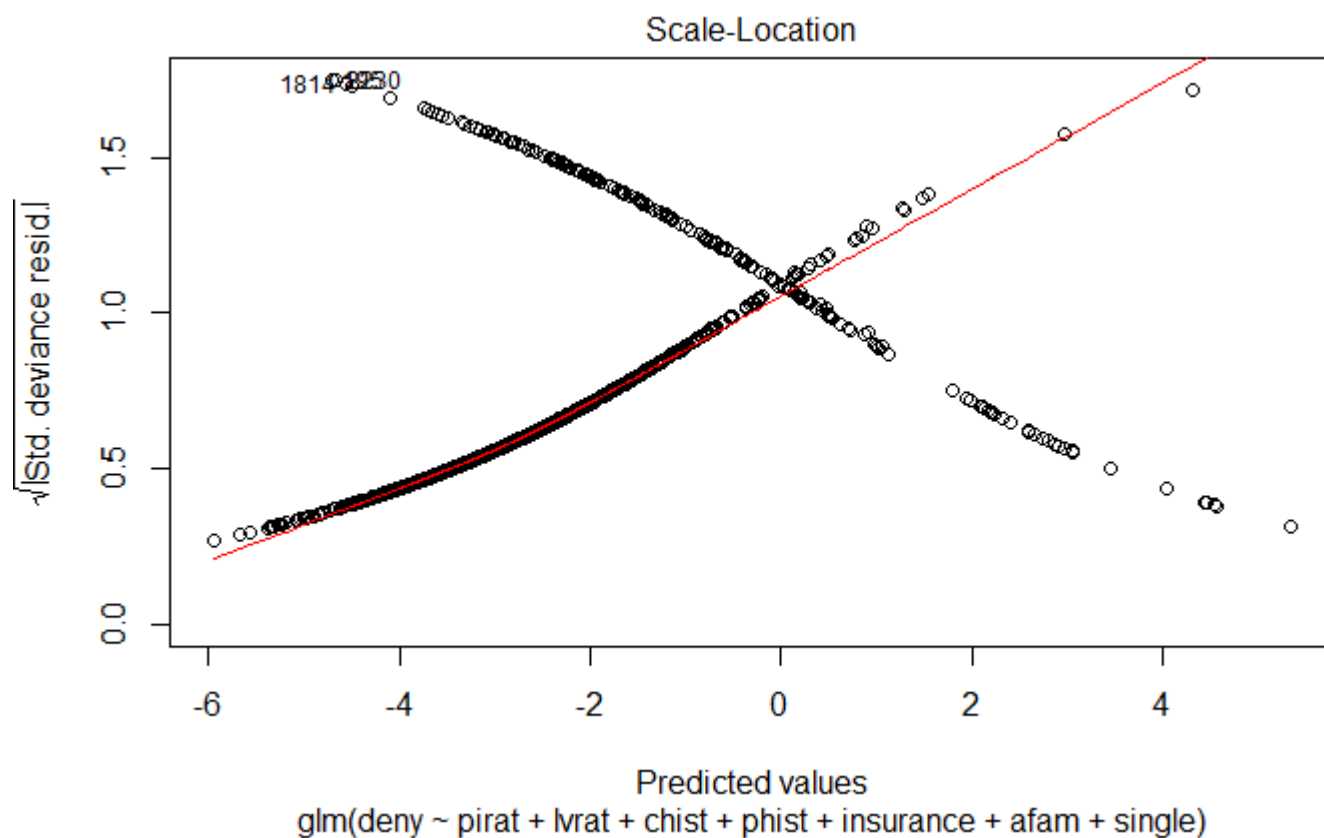
Multiple R-squared: 0.2643, Adjusted R-squared: 0.2584

F-statistic: 44.63 on 19 and 2360 DF, p-value: < 2.2e-16

Hide

plot(model3)





### Partial F-Test

Hide

```
full_model<-glm(deny~.,family=binomial(link='logit'),data=train)
model3 <- glm(deny~pirat+lvrat+chist+phist+insurance+afam+single,family=binomial(link='logit'),d
ata=train)
anova(full_model,model3)
```

### Analysis of Deviance Table

Model 1: deny ~ pirat + hirat + lvrat + chist + mhist + phist + unemp +  
selfemp + insurance + condominium + afam + single + hschool

Model 2: deny ~ pirat + lvrat + chist + phist + insurance + afam + single

	Resid. Df	Resid. Dev	Df	Deviance
1	1884	993.29		
2	1892	1013.15	-8	-19.859

### Assessing the predictive ability of the model

[Hide](#)

```
train$model3 <- predict(model2, train, type="response")
head(train)
```

deny <fctr>	pirat <dbl>	hirat <dbl>	lvrat <dbl>	chist <fctr>	mhist <fctr>	phist <fctr>	unemp <dbl>	selfemp <fctr>	
1 no	0.221	0.221	0.8000000	5	2	no	3.9	no	
2 no	0.265	0.265	0.9218750	2	2	no	3.2	no	
5 no	0.360	0.350	0.6000000	1	1	no	3.2	no	
6 no	0.240	0.170	0.5105263	1	1	no	3.9	no	
7 no	0.350	0.290	0.7466667	1	2	no	3.9	no	
8 no	0.280	0.220	0.8500000	2	2	no	1.8	no	

6 rows | 1-10 of 15 columns

[Hide](#)

```
tail(train)
```

deny <fctr>	pirat <dbl>	hirat <dbl>	lvrat <dbl>	chist <fctr>	mhist <fctr>	phist <fctr>	unemp <dbl>	selfemp <fctr>	
2374 no	0.35	0.22	0.8939394	3	2	no	3.9	no	
2375 no	0.33	0.16	0.8030303	5	1	no	3.2	no	
2376 no	0.31	0.25	0.8000000	1	1	no	3.2	yes	
2378 no	0.26	0.20	0.5267606	2	1	no	3.1	no	
2379 yes	0.32	0.26	0.7538462	6	1	yes	3.1	no	


	deny <fctr>	pirat <dbl>	hirat <dbl>	lvrat <dbl>	chist <fctr>	mhist <fctr>	phist <fctr>	unemp <dbl>	selfemp <fctr>
2380	yes	0.35	0.26	0.8135593	2	2	no	4.3	no

6 rows | 1-10 of 15 columns

Hide

```
library(gmodels)
library(ggplot2)
library(Hmisc)
library(caTools)
library(ROCR)
colAUC(train$model3, train$deny, plotROC=TRUE)
```

[,1]  
no vs. yes 0.8277181



Hide

```
predict1 <- ifelse(train$model3>0.95, 1, 0)
tab1 <- table(predicted = predict1, actual = train$deny)
tab1
```

	actual	
predicted	no	yes
0	1683	212
1	1	8

Hide

```
accuracy1<-(tab1[1,1]+tab1[2,2])/(tab1[1,1]+tab1[2,2]+tab1[1,2]+tab1[2,1])
recall1<-(tab1[2,2])/(tab1[1,2]+tab1[2,2])
precision1<-(tab1[2,2])/(tab1[2,2]+tab1[2,1])
print(c("Accuracy:",accuracy1))
```

```
[1] "Accuracy:"          "0.88813025210084"
```

Hide

```
print(c("Precision:",precision1))
```

```
[1] "Precision:"          "0.888888888888889"
```

Hide

```
print(c("Recall:",recall1))
```

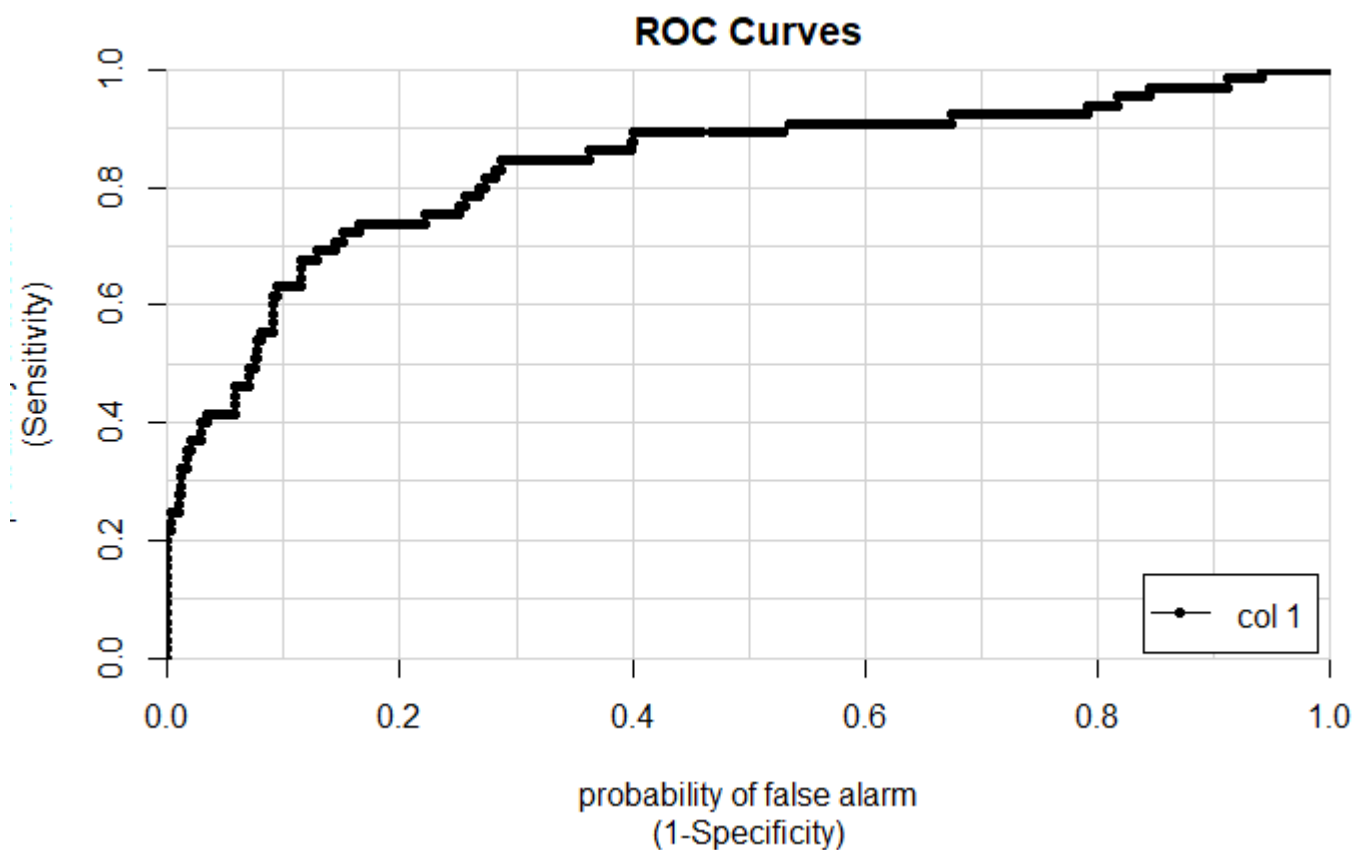
```
[1] "Recall:"             "0.0363636363636364"
```

Hide

```
test$model3 <- predict(model3, test, type='response')
colAUC(test$model3,test$deny, plotROC=TRUE)
```

```
          [,1]
no vs. yes 0.8348493
```



[Hide](#)

```
predict2 <- ifelse(test$model3>0.95, 1,0)
tab2 <- table(predicted = predict2, actual = test$deny)
tab2
```

```
      actual
predicted no yes
0      411  56
1         0   9
```

[Hide](#)

```
accuracy<-(tab2[1,1]+tab2[2,2])/(tab2[1,1]+tab2[2,2]+tab2[1,2]+tab2[2,1])
recall<-(tab2[2,2])/(tab2[1,2]+tab2[2,2])
precision<-(tab2[2,2])/(tab2[2,2]+tab2[2,1])
print(c("Accuracy:",accuracy))
```

```
[1] "Accuracy:"      "0.882352941176471"
```

[Hide](#)

```
print(c("Precision:",precision))
```

```
[1] "Precision:" "1"
```

[Hide](#)

```
print(c("Recall:",recall))
```

```
[1] "Recall:"          "0.138461538461538"
```